

Statistics is the inverse problem of probability.

Probability: we know the parameters, we figure out what the generated data are likely to look like.

Statistics: we know the data, we figure out what the parameters are likely to be.

One way of doing statistics, i.e. figuring out the unknown parameters, is to use Maximum Likelihood Estimation (MLE). (There are many other ways, e.g. OLS, Bayesian, non-parametric methods).

The idea of MLE is intuitive. Given observed data  $X$ , we find an estimate  $\hat{\theta}$  that makes  $X$  most likely to be observed, and we say that  $\hat{\theta}$  is our best guess for  $\theta$ .

For example, given 5 coin flips,  $X = \{H, T, T, T, T\}$ , with  $\theta$  is the chance of flipping head. We find the  $\hat{\theta}$  that maximizes the likelihood that we see this sequence of  $X$ . Doing the math,  $\hat{\theta} = \frac{1}{n} \sum X = 0.2$ , and that is our best guess for  $\theta$ .

While this idea makes intuitive sense, we also want to have formal proof that it's a good idea. So far we have not managed to prove anything about how good of a guess is  $\hat{\theta}$ , i.e. how close it is to  $\theta$ . So we need to learn more about the properties of MLE, by the way of learning about the expectation and variance of this estimator.

What exactly do we mean by expectation and variance of an estimator? Recall that only a random variable has an expectation and a variance. (Indeed, a number will just be ... a number). So what is random about the estimator?

Recall the coin flipping example.  $\hat{\theta} = \frac{1}{n} \sum X$ . What's random in here?  $X$  is random. Once  $X$  is observed, we got an estimate, a number. An estimator is a rule to be applied to any set of data that we will observe.

Statisticians are interested in estimators that have good properties, i.e. estimators that will give estimates that are close to the true value most of the time. Notice that I say "most of the times," because for a given dataset we have no way to know for sure how close it is to the true value.

For example, let's say the true value of  $\theta = 0.4$ .

```
theta <- 0.4
(X <- rbinom(5, size = 1, prob = theta))

## [1] 1 1 1 1 0

(MLE_estimate <- mean(X))

## [1] 0.8
```

There is no way to know for a particular dataset whether the MLE estimate is close to  $\theta$  or not. We can only know about the properties of the MLE estimate across many datasets. Some "good" properties of an estimator is unbiasedness (mean of estimator equals the true value), consistent (explain here: A sequence of  $\hat{\theta}_1, \dots, \hat{\theta}_n$  converges in probability to true value  $\limPr|\hat{\theta} - \theta| > \epsilon \leftarrow 0$ ),

efficient (an estimator that has small variance, i.e. it always stays rather close to the true value, no matter what particular dataset we get). When you go on and read methods paper on your own, when people propose new estimators, they also come up with an intuitive idea, and then have to prove that their estimator has these nice properties. So if your models are not showing significant results, you may want to search for a more efficient estimator, for example.

Side note: You've heard that MLE is not always unbiased. You wonder why we use it. It's because MLE is more efficient than OLS. (Note that OLS is BLUE, best linear UNBIASED estimator. So OLS is best among unbiased estimators, but MLE doesn't belong to that class).

Finally, talk about Fisher's information. Fisher's information quantifies how much information we have about the true value of the parameter. There are three definitions that can be shown to be algebraically equivalent:

$$I(\theta) = E \left( \left[ \frac{\partial}{\partial \theta} LL \right]^2 \right) \quad (1)$$

$$I(\theta) = -E \left( \frac{\partial^2}{\partial \theta \partial \theta} LL \right) \quad (2)$$

$$I(\theta) = Var \left( \frac{\partial}{\partial \theta} LL \right) \quad (3)$$

Most of the times, you use (2) to calculate  $I(\theta)$ , and (2) and (3) are the most intuitive.

(2) is the definition you see in class. The second derivative of the log likelihood quantifies the curvature of the log likelihood. When we zoom into the area around the estimator (which is also pretty close to the true value) we can say that it's "locally quadratic". A "blunt" support curve (one with a shallow maximum) would have a low negative expected second derivative, and thus low information; while a sharp one would have a high negative expected second derivative and thus high information.

## 1 Cobb Douglas example

$$Y = AL^\alpha K^\beta \quad (4)$$

$$\log Y = \log A + \alpha \log L + \beta \log K \quad (5)$$

We can estimate  $\alpha, \beta$ , but how would we test the hypothesis that  $H_0 : \alpha + \beta > 1$ . We can use simulation.

## 2 R stuff

```
##### Normally Distributed DVs 1 #####

library(bbmle)

## Loading required package: stats4

library(arm)

## Loading required package: MASS
## Loading required package: Matrix
## Loading required package: lme4
##
## arm (Version 1.8-6, built: 2015-7-7)
## Working directory is /home/anh/projects/ps630_lab/ps733.s16/W2

# Data from 2012 American National Election Study (~3100 non-Latino-white-identifying respon

anesdata <- na.omit(read.delim("2012 ANES_Economic Prefs.txt", header=TRUE))
summary(anesdata)
```

##	south	male	age01	unemp
##	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.00000
##	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.3333	1st Qu.:0.00000
##	Median :0.0000	Median :1.0000	Median :0.5833	Median :0.00000
##	Mean :0.3232	Mean :0.5008	Mean :0.5768	Mean :0.05109
##	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:0.8333	3rd Qu.:0.00000
##	Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.00000
##	union	income01	auth01	extrav01
##	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000
##	1st Qu.:0.0000	1st Qu.:0.2963	1st Qu.:0.2500	1st Qu.:0.3333
##	Median :0.0000	Median :0.5185	Median :0.5000	Median :0.5000
##	Mean :0.1665	Mean :0.5232	Mean :0.5643	Mean :0.5221
##	3rd Qu.:0.0000	3rd Qu.:0.7778	3rd Qu.:0.7500	3rd Qu.:0.6667
##	Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.0000
##	agree01	consc01	stable01	openness01
##	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000
##	1st Qu.:0.5833	1st Qu.:0.6667	1st Qu.:0.5000	1st Qu.:0.5000
##	Median :0.6667	Median :0.8333	Median :0.6667	Median :0.6667
##	Mean :0.6971	Mean :0.7818	Mean :0.6557	Mean :0.6320
##	3rd Qu.:0.8333	3rd Qu.:0.9167	3rd Qu.:0.8333	3rd Qu.:0.7500
##	Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.0000
##	econ01	educ1	educ2	educ3
##	Min. :0.0000	Min. :0.00000	Min. :0.000	Min. :0.0000
##	1st Qu.:0.2333	1st Qu.:0.00000	1st Qu.:0.000	1st Qu.:0.0000
##	Median :0.4333	Median :0.00000	Median :0.000	Median :0.0000
##	Mean :0.4306	Mean :0.06465	Mean :0.245	Mean :0.3198
##	3rd Qu.:0.6000	3rd Qu.:0.00000	3rd Qu.:0.000	3rd Qu.:1.0000

```

## Max. :1.0000 Max. :1.00000 Max. :1.000 Max. :1.0000
## educ4 educ5
## Min. :0.0000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.0000 Median :0.0000
## Mean :0.2208 Mean :0.1498
## 3rd Qu.:0.0000 3rd Qu.:0.0000
## Max. :1.0000 Max. :1.0000

### Simple linear regression of social welfare support on income (coded 0-1) and union membership

X <- array(NA, c(length(anesdata$secon01),3)) #Initialize the design matrix
X[,1] <- 1 #column of 1s for constant
X[,2] <- anesdata$income01 #household income, coded from 0-1
X[,3] <- anesdata$union #dichotomous indicator for family member union membership
y <- anesdata$secon01 #additive scale of several social welfare policy items, coded from 0-1

# Define the LL function in general terms (this code can be used for any # of predictors):

LL_normreg = function(params, y, X){
  B = matrix(NA, nrow = length(params) - 1, ncol = 1)
  B[,1] = params[-length(params)]
  sigma = params[length(params)]
  minusll = -sum(dnorm(y, X %*% B, sigma, log=T))
  return(minusll)
}

# Declare the names of the parameters (from B0 to B[# of predictors], and sigma):

parnames(LL_normreg) <- c("B0", "B1", "B2", "sigma")

# Fit the model using mle2 ('vecpar=TRUE' tells mle2 that the first argument passed to the
# LL function is a vector of all parameters with names declared in 'parnames' above and in

fit <- mle2(LL_normreg, start = c(B0 = mean(y), B1 = 0, B2 = 0, sigma = sd(y)),
  data=list(y=y,X=X), vecpar = TRUE, control=list(maxit=5000))

## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced

```

```
## Warning in dnorm(y, X %*% B, sigma, log = T): NaNs produced

summary(fit)

## Maximum likelihood estimation
##
## Call:
## mle2(minuslogl = LL_normreg, start = c(B0 = mean(y), B1 = 0,
##     B2 = 0, sigma = sd(y)), data = list(y = y, X = X), vecpar = TRUE,
##     control = list(maxit = 5000))
##
## Coefficients:
##           Estimate Std. Error z value      Pr(z)
## B0      0.4686993   0.0085633  54.7333 < 2.2e-16 ***
## B1     -0.0931828   0.0143522  -6.4926 8.439e-11 ***
## B2      0.0639140   0.0111607   5.7267 1.024e-08 ***
## sigma   0.2322673   0.0029167  79.6349 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## -2 log L: -259.6571

str(summary(fit))

## Formal class 'summary.mle2' [package "bbmle"] with 3 slots
##   ..@ call   : language mle2(minuslogl = LL_normreg, start = c(B0 = mean(y), B1 = 0, B2 =
##   ..@ coef    : num [1:4, 1:4] 0.4687 -0.09318 0.06391 0.23227 0.00856 ...
##   .. ..- attr(*, "dimnames")=List of 2
##   .. .. ..$ : chr [1:4] "B0" "B1" "B2" "sigma"
##   .. .. ..$ : chr [1:4] "Estimate" "Std. Error" "z value" "Pr(z)"
##   ..@ m2logL: num -260

summary(fit)@m2logL

## [1] -259.6571
```