

Pol Sci 630: Problem Set 4 - Regression Model Estimation

Prepared by: Anh Le (anh.le@duke.edu)

Due Date: Tuesday, September 22nd, 2015, 10 AM (Beginning of Class)

Note 1: It is absolutely essential that you show all your work, including intermediary steps, and comment on your R code to earn full credit.

Note 2: Please use a *single* PDF file created through knitr to submit your answers. knitr allows you to combine R code and L^AT_EX code in one document, meaning that you can include both the answers to R programming and math problems. Also submit the source code that generates the PDF file (i.e. .Rnw file)

Note 3: Make sure that the PDF files you submit do not include any references to your identity. The grading will happen anonymously. You can submit your answer at the following website: <http://ps630-f15.herokuapp.com/>

1. Create a data frame (4 points)

a)

First, `set.seed(2)`. Then, create a data frame with 1000 rows and 3 variables as follows:

1. `var_norm`: a normal variable with mean = 5, sd = 10
2. `var_binom`: a binomial variable with number of trial = 10, probability of success = 0.5
3. `var_poisson`: a Poisson variable with $\lambda = 4$

(Recall how to generate random sample from various distributions from previous labs.)

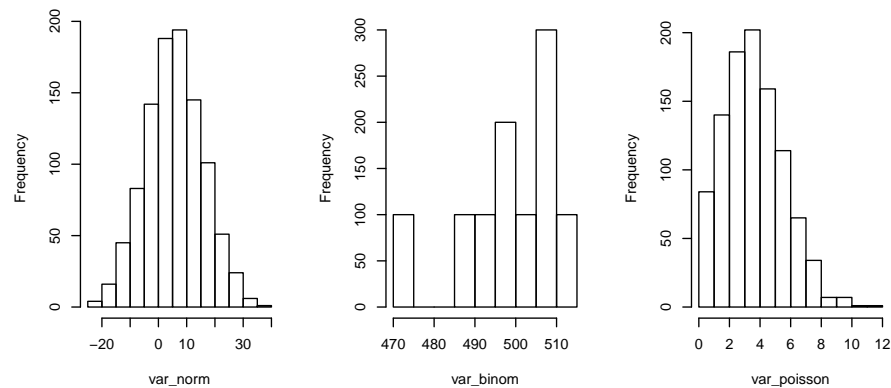
b)

Plot the histograms of the three variables, arranging them nicely (with `fig.width()`, `fig.height()`, `par(mfrow)` as you see fit). Brownie point if you plot using a for loop instead of writing `hist` three times.

Solution

```
# Create the data frame
set.seed(2)
my_dataframe <- data.frame(var_norm = rnorm(1000, mean = 5, sd = 10),
                           var_binom = rbinom(1000, n = 10, prob = 0.5),
                           var_poisson = rpois(1000, lambda = 4))

# Plot the histogram (nicely)
par(mfrow = c(1, 3))
for (i in 1:3) {
  hist(my_dataframe[, i],
       xlab = colnames(my_dataframe)[i], main = NULL)
}
```



2. Subset data frame (4 points)

a)

Download the following data from WDI and clean it as follows. Briefly comment on what each command does.

```
library(WDI)

## Loading required package: RJSONIO

d_wdi <- WDI(indicator = c("NY.GDP.PCAP.CD", "SP.DYN.IMRT.IN", "SH.MED.PHYS.ZS"),
             start = 2005, end = 2010, extra = TRUE)
d_wdi <- d_wdi[d_wdi$region != "Aggregates",
              c("country", "year", "NY.GDP.PCAP.CD", "SP.DYN.IMRT.IN", "SH.MED.PHYS.ZS")]
colnames(d_wdi)[3:5] <- c('gdppc', 'infant_mortality', 'number_of_physician')
d_wdi <- na.omit(d_wdi)
```

infant_mortality: number of mortality per 1000 live births
number_of_physician: number of physician per 1000 people

b)

Use subsetting techniques to do the following:

1. Show the GDP per capita of Brazil across years
2. Show the country-years where infant mortality > 100 per 1000 live birth
3. Show the country-years where GDP per capita is above average
4. Show the country-years where GDP per capita is above average, but number of physician is below average

Solution

```
library(WDI)

# Download data from WDI, specifying the indicators and start / end year
d_wdi <- WDI(indicator = c("NY.GDP.PCAP.CD", "SP.DYN.IMRT.IN", "SH.MED.PHYS.ZS"),
             start = 2008, end = 2010, extra = TRUE)

# Remove aggregates rows, selecting wanted columns by name
d_wdi <- d_wdi[d_wdi$region != "Aggregates",
              c("country", "year", "NY.GDP.PCAP.CD", "SP.DYN.IMRT.IN", "SH.MED.PHYS.ZS")]

# Rename some of the columns
colnames(d_wdi)[3:5] <- c('gdppc', 'infant_mortality', 'number_of_physician')

# Remove all rows that have missing data
d_wdi <- na.omit(d_wdi)

# 1. Show the GDP per capita of Brazil across years
d_wdi[d_wdi$country == "Brazil", c("country", "year", "gdppc")]

##      country year      gdppc
## 94  Brazil 2008  8836.914
## 95  Brazil 2010 11318.057

# 2. Show the country-years where infant mortality > 100 per 1000 live birth
d_wdi[d_wdi$infant_mortality > 100, c("country", "year", "infant_mortality")]

##              country year infant_mortality
## 34              Angola 2009             112.2
## 120 Central African Republic 2009             103.6
## 562              Sierra Leone 2010             107.0
## 563              Sierra Leone 2008             116.2
```

```

# 3. Show the country-years where GDP per capita is above average
d_wdi[d_wdi$gdppc > mean(d_wdi$gdppc), c("country", "year", "gdppc")]

##              country year      gdppc
## 16              Andorra 2010 42952.72
## 17              Andorra 2009 46401.09
## 20  United Arab Emirates 2010 33885.93
## 43              Austria 2010 46593.39
## 48              Australia 2010 51801.05
## 62              Barbados 2010 15854.36
## 67              Belgium 2010 44360.90
## 69              Belgium 2008 48561.36
## 76              Bahrain 2008 23037.64
## 77              Bahrain 2010 20545.75
## 88  Brunei Darussalam 2008 37094.86
## 89  Brunei Darussalam 2010 30882.40
## 99              Bahamas, The 2008 23674.14
## 113             Canada 2008 46400.44
## 114             Canada 2010 47463.63
## 124             Switzerland 2010 74277.12
## 154             Cyprus 2010 30438.90
## 155             Cyprus 2008 34950.35
## 157  Czech Republic 2008 22649.38
## 158  Czech Republic 2010 19763.96
## 160             Germany 2008 45632.84
## 162             Germany 2010 41725.85
## 167             Denmark 2009 57895.50
## 168             Denmark 2010 57647.67
## 182             Estonia 2008 18087.68
## 190              Spain 2010 30737.83
## 202             Finland 2009 47107.16
## 203             Finland 2010 46205.17
## 204             Finland 2008 53401.31
## 214             France 2008 45415.81
## 215             France 2010 40708.50
## 222  United Kingdom 2010 38362.22
## 245             Greece 2010 26863.01
## 246             Greece 2008 31700.49
## 267             Croatia 2008 15887.42
## 273             Hungary 2008 15598.32
## 278             Ireland 2008 60968.84
## 279             Ireland 2010 47903.68
## 280             Israel 2010 30551.12
## 295             Iceland 2008 55446.76
## 297             Iceland 2010 41695.89
## 298              Italy 2009 36995.11

```

```

## 299          Italy 2010 35877.87
## 300          Italy 2008 40659.67
## 310          Japan 2010 42909.23
## 312          Japan 2008 37865.62
## 334      Korea, Rep. 2008 20474.89
## 336      Korea, Rep. 2010 22151.21
## 337          Kuwait 2010 38580.41
## 338          Kuwait 2008 54540.42
## 339          Kuwait 2009 37158.42
## 367      Lithuania 2008 14961.72
## 371      Luxembourg 2010 102863.10
## 423          Malta 2010 19694.08
## 458      Netherlands 2010 50341.25
## 459      Netherlands 2008 56628.75
## 460          Norway 2009 80017.78
## 461          Norway 2010 87646.27
## 462          Norway 2008 96880.51
## 467      New Zealand 2010 33394.07
## 472          Oman 2010 20922.66
## 474          Oman 2008 23483.63
## 504      Portugal 2010 22539.99
## 512          Qatar 2010 71510.19
## 538      Saudi Arabia 2008 19714.40
## 539      Saudi Arabia 2010 19326.58
## 540      Saudi Arabia 2009 16013.28
## 550          Sweden 2008 55746.84
## 551          Sweden 2010 52076.43
## 552          Sweden 2009 46207.06
## 553      Singapore 2010 46569.69
## 556      Slovenia 2008 27501.82
## 558      Slovenia 2010 23417.64
## 560      Slovak Republic 2010 16509.90
## 626      Trinidad and Tobago 2010 15494.70
## 643      United States 2010 48374.06
## 644      United States 2009 47001.56

# 4. Show the country-years where GDP per capita is above average,
# but number of physician is below average
d_wdi[d_wdi$gdppc > mean(d_wdi$gdppc) &
      d_wdi$number_of_physician < mean(d_wdi$number_of_physician),
      c("country", "year", "gdppc")]

##          country year    gdppc
## 76          Bahrain 2008 23037.64
## 77          Bahrain 2010 20545.75
## 88      Brunei Darussalam 2008 37094.86

```

```
## 89      Brunei Darussalam 2010 30882.40
## 538      Saudi Arabia 2008 19714.40
## 539      Saudi Arabia 2010 19326.58
## 540      Saudi Arabia 2009 16013.28
## 626 Trinidad and Tobago 2010 15494.70
```

3. Build linear model (4 points)

a)

Download 2 variables of interest and build a linear model of their relationship using `lm()`. Show the `summary()` of results

b)

Show the result with `stargazer`, customizing:

- The labels of the independent variables (i.e. the covariate)
- The label of the dependent variable
- Make the model name (i.e. OLS) show up

Hint: The options to do those things are in `help(stargazer)`. I have worded the task in a way that should help you find the relevant options.

Solution

```
m1 <- lm(infant_mortality ~ gdppc, data = d_wdi)
summary(m1)

##
## Call:
## lm(formula = infant_mortality ~ gdppc, data = d_wdi)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -28.743 -17.413  -5.357   11.922   78.783
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.775e+01  1.713e+00   22.04  <2e-16 ***
## gdppc        -7.406e-04  6.908e-05  -10.72  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.68 on 249 degrees of freedom
```

```
## Multiple R-squared:  0.3158, Adjusted R-squared:  0.3131
## F-statistic: 114.9 on 1 and 249 DF,  p-value: < 2.2e-16
```

```
library(stargazer)

##
## Please cite as:
##
## Hlavac, Marek (2014). stargazer: LaTeX code and ASCII text for
## well-formatted regression and summary statistics tables.
## R package version 5.1. http://CRAN.R-project.org/package=stargazer

stargazer(m1,
  covariate.labels = c("GDP per capita"),
  dep.var.labels = c("Infant Mortality (per 1000 births)"),
  model.names = TRUE)
```

Table 1:

<i>Dependent variable:</i>	
Infant Mortality (per 1000 births)	
<i>OLS</i>	
GDP per capita	−0.001*** (0.0001)
Constant	37.753*** (1.713)
Observations	251
R ²	0.316
Adjusted R ²	0.313
Residual Std. Error	21.678 (df = 249)
F Statistic	114.948*** (df = 1; 249)
<i>Note:</i> *p<0.1; **p<0.05; ***p<0.01	

4. Calculate sum of squares and RMSE (4 points)

1. Extract the residuals and predicted values (fitted values) from the model object (from Question 3)
2. Calculate three “sum of squares” (TSS, RegSS, RSS)

3. Calculate the root mean square error and compare with R. (In R and stargazer, RMSE is called “Residual standard error”.)

Note: the data you feed to `lm()` may have missing data, so R has to modify the data a little before using it. To extract the data that are actually used by `lm()`, use `my_model$model`. Use this data to calculate \bar{y} in the sum of squares.

Solution

```
res <- m1$residuals # Residuals
pred <- m1$fitted.values # Predicted values
y <- m1$model$infant_mortality # Data of Y that is used by lm()

# Calculate 3 sum of squares
TSS <- sum( (y - mean(y)) ** 2 )
RegSS <- sum( (pred - mean(y)) ** 2 )
RSS <- sum( res ** 2 )

# Calculate root mean square error
N <- nrow(d_wdi)
k <- 1 # We only have 1 predictor, which is log_gdppc
rmse <- sqrt(RSS / (N - k - 1))
```

The calculated root mean square error is 21.6783991, the same as reported by R in `summary(m1)`.