

Tutorial 4: Regression Model Estimation

Anh Le (anh.le@duke.edu)

September 18, 2015

Agenda

1. Create data frames
 2. Subset data frames
 3. Estimate a linear model with `lm()`
 4. Tips and tricks
- Prefix your objects in R (and related TAB tricks, i.e. arguments within function, variables within a data frame)
 - `fig.height()`, `fig.width()`
 - Code length ≤ 80 (RStudio > Tools > Options > Code)

1. Create data frames

```
my_dataframe <- data.frame(var1 = c(11, 12, 13),
                           var2 = c(21, 22, 23),
                           var3 = c("a", "b", "c"))
my_dataframe
```

```
##   var1 var2 var3
## 1   11   21    a
## 2   12   22    b
## 3   13   23    c
```

2. Subset data frames

All subsetting can be done with the following construct: `my_dataframe[?1 , ?2]`

The first question mark (?1) refers to which rows we want. The second question mark (?2) refers to which columns we want.

How to indicate to R which rows / columns we want? Multiple ways:

1. Use rows / columns index

```
my_dataframe[1, 2]
```

```
## [1] 21
```

```
my_dataframe[1:2, 2]
```

```
## [1] 21 22
```

```
my_dataframe[1:2, ]
```

```
##   var1 var2 var3
## 1   11   21    a
## 2   12   22    b
```

Rapid fire quiz

```
my_dataframe[2:3, ]
my_dataframe[ , 1:2]
my_dataframe[1:2, 2:3]

my_dataframe[c(1, 3), ]
my_dataframe[c(1, 3, 2), ]
```

2. Use rows / columns name

```
my_dataframe[ , "var2"]
```

```
## [1] 21 22 23
```

Rapid fire quiz:

```
my_dataframe[ , c("var1", "var3")]
my_dataframe[c(2, 3), c("var1", "var3")]
```

3. Use a condition

```
my_dataframe[c(TRUE, TRUE, FALSE), ]
```

```
##   var1 var2 var3
## 1   11   21    a
## 2   12   22    b
```

```
my_dataframe[, c(TRUE, FALSE, TRUE)]
```

```
##   var1 var3
## 1   11    a
## 2   12    b
## 3   13    c
```

Of course this is not tenable for a large data frame. So we have this very useful trick:

```
my_dataframe[my_dataframe$var1 < 13, ]
```

```
##   var1 var2 var3
## 1   11   21    a
## 2   12   22    b
```

This works because `my_dataframe$var1 < 13` actually returns `c(TRUE, TRUE, FALSE)` (vectorized operation in the wild!). Indeed:

```
my_dataframe$var1 < 13
```

```
## [1]  TRUE  TRUE FALSE
```

Rapid fire quiz:

```
my_dataframe[my_dataframe$var2 == 22, ]
my_dataframe[my_dataframe$var2 == 25, ]
```

4. Use a combination of condition

```
my_dataframe[my_dataframe$var1 > 10 & my_dataframe$var2 > 21, ]
```

```
##   var1 var2 var3
## 2   12   22    b
## 3   13   23    c
```

```
my_dataframe[my_dataframe$var1 > 10 | my_dataframe$var2 > 21, ]
```

```
##   var1 var2 var3
## 1   11   21    a
## 2   12   22    b
## 3   13   23    c
```

3. Estimate a linear model with `lm()`

In this section, I'll demo a (simplified) pipeline of steps in doing regression analysis with real data.

Download and clean data

```
library(WDI)
```

```
## Loading required package: RJSONIO
```

```
d_2010 <- WDI(indicator = c("NY.GDP.PCAP.CD", "SP.DYN.IMRT.IN", "SH.MED.PHYS.ZS"),
              start = 2010, end = 2010, extra = TRUE)
# d_2010[d_2010$region != "Aggregates", ]
```

There are a lot of unwanted columns. What if I just want `country`, `year`, and the three variables of interest (NY.GDP.PCAP.CD, SP.DYN.IMRT.IN, SH.MED.PHYS.ZS)? (Hint: subsetting)

```
d_2010 <- d_2010[d_2010$region != "Aggregates",
                 c("country", "year", "NY.GDP.PCAP.CD", "SP.DYN.IMRT.IN", "SH.MED.PHYS.ZS")]
```

Rename columns:

```
colnames(d_2010)
```

```
## [1] "country"      "year"          "NY.GDP.PCAP.CD" "SP.DYN.IMRT.IN"
## [5] "SH.MED.PHYS.ZS"
```

```
colnames(d_2010)[3:5] <- c('gdppc', 'infant_mortality', 'number_of_physician')
colnames(d_2010)
```

```
## [1] "country"      "year"          "gdppc"
## [4] "infant_mortality" "number_of_physician"
```

Log gdp per capita

```
d_2010$log_gdppc <- log(d_2010$gdppc)
```

Remove missing data

```
d_2010 <- na.omit(d_2010)
```

Build a linear model

```
lm(infant_mortality ~ log_gdppc, data = d_2010)
```

```
##
## Call:
## lm(formula = infant_mortality ~ log_gdppc, data = d_2010)
##
## Coefficients:
## (Intercept)      log_gdppc
##      134.35         -12.76
```

```
m1 <- lm(infant_mortality ~ log_gdppc, data = d_2010)
summary(m1)
```

```
##
## Call:
## lm(formula = infant_mortality ~ log_gdppc, data = d_2010)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -25.652  -8.956  -0.659   6.673  50.489
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  134.3454     6.2980   21.33  <2e-16 ***
## log_gdppc    -12.7574     0.7205  -17.70  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.05 on 141 degrees of freedom
## Multiple R-squared:  0.6898, Adjusted R-squared:  0.6876
## F-statistic: 313.5 on 1 and 141 DF,  p-value: < 2.2e-16
```

```
m2 <- lm(infant_mortality ~ log_gdppc + number_of_physician, data = d_2010)
summary(m2)
```

```
##
## Call:
## lm(formula = infant_mortality ~ log_gdppc + number_of_physician,
##      data = d_2010)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -27.299  -8.117  -1.109   5.916  50.530
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    121.3864     7.4036  16.396  < 2e-16 ***
## log_gdppc      -10.6297     0.9794 -10.853  < 2e-16 ***
## number_of_physician -2.8746     0.9262  -3.103  0.00231 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.67 on 140 degrees of freedom
## Multiple R-squared:  0.7097, Adjusted R-squared:  0.7056
## F-statistic: 171.2 on 2 and 140 DF,  p-value: < 2.2e-16
```

Extract result from the model

`str()` (stands for structure) is used to look into the structure of an object in R, see what it contains.

```
str(m1)
```

```
## List of 12
## $ coefficients : Named num [1:2] 134.3 -12.8
## ..- attr(*, "names")= chr [1:2] "(Intercept)" "log_gdppc"
```

```

## $ residuals      : Named num [1:143] 3.22 6.19 21.71 -13.44 -15.59 ...
##   ..- attr(*, "names")= chr [1:143] "6" "7" "8" "10" ...
## $ effects        : Named num [1:143] -293.3 231.1 22.2 -13.6 -15.6 ...
##   ..- attr(*, "names")= chr [1:143] "(Intercept)" "log_gdppc" "" "" ...
## $ rank           : int 2
## $ fitted.values: Named num [1:143] -0.725 1.106 53.393 28.237 31.685 ...
##   ..- attr(*, "names")= chr [1:143] "6" "7" "8" "10" ...
## $ assign         : int [1:2] 0 1
## $ qr             :List of 5
##   ..$ qr        : num [1:143, 1:2] -11.9583 0.0836 0.0836 0.0836 0.0836 ...
##   .. ..- attr(*, "dimnames")=List of 2
##   .. .. ..$ : chr [1:143] "6" "7" "8" "10" ...
##   .. .. ..$ : chr [1:2] "(Intercept)" "log_gdppc"
##   .. ..- attr(*, "assign")= int [1:2] 0 1
##   ..$ qraux: num [1:2] 1.08 1.09
##   ..$ pivot: int [1:2] 1 2
##   ..$ tol   : num 1e-07
##   ..$ rank  : int 2
##   ..- attr(*, "class")= chr "qr"
## $ df.residual    : int 141
## $ xlevels        : Named list()
## $ call           : language lm(formula = infant_mortality ~ log_gdppc, data = d_2010)
## $ terms          :Classes 'terms', 'formula' length 3 infant_mortality ~ log_gdppc
##   .. ..- attr(*, "variables")= language list(infant_mortality, log_gdppc)
##   .. ..- attr(*, "factors")= int [1:2, 1] 0 1
##   .. .. ..- attr(*, "dimnames")=List of 2
##   .. .. .. ..$ : chr [1:2] "infant_mortality" "log_gdppc"
##   .. .. .. ..$ : chr "log_gdppc"
##   .. ..- attr(*, "term.labels")= chr "log_gdppc"
##   .. ..- attr(*, "order")= int 1
##   .. ..- attr(*, "intercept")= int 1
##   .. ..- attr(*, "response")= int 1
##   .. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
##   .. ..- attr(*, "predvars")= language list(infant_mortality, log_gdppc)
##   .. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
##   .. .. ..- attr(*, "names")= chr [1:2] "infant_mortality" "log_gdppc"
## $ model          :'data.frame': 143 obs. of 2 variables:
##   ..$ infant_mortality: num [1:143] 2.5 7.3 75.1 14.8 16.1 13 3.6 4.1 33.9 6.4 ...
##   ..$ log_gdppc       : num [1:143] 10.59 10.44 6.35 8.32 8.05 ...
##   ..- attr(*, "terms")=Classes 'terms', 'formula' length 3 infant_mortality ~ log_gdppc
##   .. .. ..- attr(*, "variables")= language list(infant_mortality, log_gdppc)
##   .. .. ..- attr(*, "factors")= int [1:2, 1] 0 1
##   .. .. .. ..- attr(*, "dimnames")=List of 2
##   .. .. .. .. ..$ : chr [1:2] "infant_mortality" "log_gdppc"
##   .. .. .. .. ..$ : chr "log_gdppc"
##   .. .. ..- attr(*, "term.labels")= chr "log_gdppc"
##   .. .. ..- attr(*, "order")= int 1
##   .. .. ..- attr(*, "intercept")= int 1
##   .. .. ..- attr(*, "response")= int 1
##   .. .. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
##   .. .. ..- attr(*, "predvars")= language list(infant_mortality, log_gdppc)
##   .. .. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
##   .. .. .. ..- attr(*, "names")= chr [1:2] "infant_mortality" "log_gdppc"
## - attr(*, "class")= chr "lm"

```

You can extract the coefficients

```
m1$coefficients
```

```
## (Intercept)    log_gdppc  
##      134.3454     -12.7574
```

```
m1$coefficients['(Intercept)']
```

```
## (Intercept)  
##      134.3454
```

```
m1$coefficients['log_gdppc']
```

```
## log_gdppc  
##     -12.7574
```


You can also generate predicted / fitted values:

```
d_2010$pred_infant_mortality1 <- predict(m1)  
d_2010$pred_infant_mortality2 <- m1$coefficients['(Intercept)'] + m1$coefficients['log_gdppc'] * d_2010$log_gdppc
```


Now we can use them for other things, e.g plotting

```
par(mfrow = c(1, 2))  
plot(infant_mortality ~ log_gdppc, data = d_2010, main = "Plot predicted values-method 1")  
abline(a = m1$coefficients['(Intercept)'], b = m1$coefficients['log_gdppc'])  
points(d_2010$log_gdppc, d_2010$pred_infant_mortality1, col = 'red')  
  
plot(infant_mortality ~ log_gdppc, data = d_2010, main = "Plot predicted values-method 2")  
abline(a = m1$coefficients['(Intercept)'], b = m1$coefficients['log_gdppc'])  
points(d_2010$log_gdppc, d_2010$pred_infant_mortality2, col = 'red')
```

Plot predicted values–method 1



Plot predicted values–method 2



Report the model in a nice, journal-ready format

The `stargazer` library takes your model objects and generates tables in LaTeX. This package has a lot of customizing options, which you'll explore in the homework.

```
library(stargazer)
```

```
##
## Please cite as:
##
## Hlavac, Marek (2014). stargazer: LaTeX code and ASCII text for well-formatted regression and summary
## R package version 5.1. http://CRAN.R-project.org/package=stargazer
```

```
# LaTeX code that you can copy paste into LaTeX
stargazer(m1, m2)
```

```
##
## % Table created by stargazer v.5.1 by Marek Hlavac, Harvard University. E-mail: hlavac at fas.harvard
## % Date and time: Fri, Sep 18, 2015 - 02:58:23 PM
## \begin{table}[!htbp] \centering
##   \caption{}
##   \label{}
##   \begin{tabular}{@{\extracolsep{5pt}}lcc}
##     \hline
##     \hline \hline
##     & \multicolumn{2}{c}{\textit{Dependent variable:}} & \\
##     \cline{2-3}
```



```
## \[-1.8ex] & \multicolumn{2}{c}{infant\_mortality} \\
## \[-1.8ex] & (1) & (2)\\
## \hline \[-1.8ex]
## log\_gdppc & $-12.757^{***}$ & $-10.630^{***}$ \\
## & (0.721) & (0.979) \\
## & & \\
## number\_of\_physician & & $-2.875^{***}$ \\
## & & (0.926) \\
## & & \\
## Constant & 134.345^{***} & 121.386^{***} \\
## & (6.298) & (7.404) \\
## & & \\
## \hline \[-1.8ex]
## Observations & 143 & 143 \\
## R^2 & 0.690 & 0.710 \\
## Adjusted R^2 & 0.688 & 0.706 \\
## Residual Std. Error & 13.051 (df = 141) & 12.669 (df = 140) \\
## F Statistic & 313.482^{***} (df = 1; 141) & 171.152^{***} (df = 2; 140) \\
## \hline
## \hline \[-1.8ex]
## \textit{Note:} & \multicolumn{2}{r}{^{*}p<$0.1; ^{* *}p<$0.05; ^{***}p<$0.01} \\
## \end{tabular}
## \end{table}
```

```
# If using knitr, use the option results='asis'
stargazer(m1, m2)
```

% Table created by stargazer v.5.1 by Marek Hlavac, Harvard University. E-mail: hlavac at fas.harvard.edu
 % Date and time: Fri, Sep 18, 2015 - 02:58:23 PM

| Table 1: | | |
|-----------------------------|----------------------------|--------------------------|
| | <i>Dependent variable:</i> | |
| | infant_mortality | |
| | (1) | (2) |
| log_gdppc | -12.757*** (0.721) | -10.630*** (0.979) |
| number_of_physician | | -2.875*** (0.926) |
| Constant | 134.345*** (6.298) | 121.386*** (7.404) |
| Observations | 143 | 143 |
| R ² | 0.690 | 0.710 |
| Adjusted R ² | 0.688 | 0.706 |
| Residual Std. Error | 13.051 (df = 141) | 12.669 (df = 140) |
| F Statistic | 313.482*** (df = 1; 141) | 171.152*** (df = 2; 140) |
| <i>Note:</i> | | |
| *p<0.1; **p<0.05; ***p<0.01 | | |